Table 1: Near-Capacity vs. Slack-Capacity: Treated-Only vs. Including Never-Treated Controls

	Spare capacity	**Occupancy rate**
Panel A: Treated-only (baseline specification)		
Slack-capacity: *post*	-0.0387^{***}	3.8165^{***}
	(0.0048)	(0.4794)
Difference (Near − Slack)	0.0292^{***}	-2.8940^{***}
	(0.0061)	(0.6107)
Near-capacity effect	-0.0095^{*}	0.9224^{*}
	(0.0053)	(0.5294)
Observations	117,863	
Panel B: Including never-treated facilities as controls		
Slack-capacity: *post*	-0.0397^{***}	3.9573^{***}
	(0.0043)	(0.4369)
Difference (Near − Slack)	0.0279^{***}	-2.7637^{***}
	(0.0062)	(0.6170)
Near-capacity effect	-0.0117^{**}	1.1935^{***}
	(0.0046)	(0.4530)
Observations	564,605	

Notes: Combined median spare capacity used for the near/slack split: 0.167. Near-capacity = baseline spare capacity $\leq$ median; Slack-capacity = above.

Panel A restricts the sample to ever-treated (CHOW) facilities, classified by their own pre-acquisition average spare capacity (event_time $\in [-12, -4]$). This is the specification already sent to advisors.

Panel B adds never-treated facilities to the estimation sample as controls. Because never-treated facilities have no acquisition event, they cannot be classified on a pre-acquisition window – they are instead classified using a FIXED CALENDAR baseline (2017 Q1 average spare capacity), the same convention used elsewhere in this project for time-invariant baseline traits (e.g. chain status). *post* is always 0 for never-treated facilities, so they do not mechanically enter the post or interaction coefficients directly; including them changes what the calendar fixed effects (and therefore *post*) are estimated relative to.

All specifications include facility and calendar-month fixed effects and the standard control set (excluding occupancy rate/spare capacity from each other's controls). Anticipation window excluded. Standard errors clustered two ways by facility and calendar month.

Statistical significance: $^{***}p < 0.01$, $^{**}p < 0.05$, $^{*}p < 0.10$.